

MAANAS TYAGI

+91-9785307374 | maanastyagi2006@gmail.com | maanastyagi.com
linkedin.com/in/maanas-tyagi-1b955b23b | github.com/maanas1234

PROFESSIONAL SUMMARY

AI/ML engineer with hands-on experience building production AI systems, automation pipelines, and evaluated retrieval-augmented generation (RAG) pipelines. Proficient in machine learning, deep learning, NLP, and multi-agent agentic AI architectures. Author of one preprint and one paper currently under review on large language model (LLM) behavior.

SELECTED IMPACT

- Demonstrated through a 33-query, RAGAS-scored ablation study that semantic-only retrieval outperforms hybrid BM25+semantic retrieval on every metric (Recall@5: 0.848 vs. 0.818, MRR: 0.703 vs. 0.610), overturning the assumed hybrid-retrieval default
- Replaced a \$3,000 licensed reconciliation tool with a custom Python automation pipeline at Genpact, reducing 10,000-row file processing time from over 30 minutes to under 5 minutes for Fortune 500 clients; now actively used by tens of colleagues to automate their daily reconciliation work
- Contributed 5 merged pull requests to Hugging Face repositories (tau: #314, #290, #289; Transformers: #45150), each reviewed and accepted by project maintainers
- Designed and built end-to-end AI systems at Sstudize, including system architecture and the application layer for Student GPT, along with a token economics model for sustainable per-user LLM cost structure
- Co-built TrenTorch to 130+ GitHub stars in 2 days of launch — one of the fastest-growing open-source educational ML repositories on GitHub, with 10+ forks

PROFESSIONAL EXPERIENCE

LLM Engineer

July 2026 – Present

Sstudize · Remote

- Architect and own all AI systems for the department, including system design and the application layer for Student GPT, an AI-powered platform for students
- Developed a token economics model to define sustainable per-user LLM cost and usage structure for the product

AI & Analytics Intern (*Forward Deployed Engineer Level Work*)

May 2026 – June 2026

Genpact · Jaipur, India

- Replaced a \$3,000 licensed reconciliation tool with a custom Python automation pipeline, reducing processing time for 10,000-row files from over 30 minutes to under 5 minutes; now actively used by tens of colleagues for daily reconciliation work
- Built purchase, invoice, and GRNI reconciliation engines for Fortune 500 clients, converting manual pivot-table workflows into single-file uploads with instant, fully cross-matched output

PUBLICATIONS

- Instruction Fidelity Decay in Small Language Model Agents: Semantic Recall Dissociates from Format Compliance Under Semantic Interference** (Under Review)
- Front-Loading Bias in Large Language Models: A Multi-Model Empirical Analysis** (ResearchGate Preprint) – doi.org/10.13140/RG.2.2.32504.92164

PROJECTS

RAG Eval System | LangChain · Chroma · RAGAS · Python

July 2026

github.com/maanas1234/RAG-Eval-System

- Constructed a retrieval and generation evaluation harness using Precision/Recall/MRR and RAGAS faithfulness/relevancy metrics, tested against 33 hand-labeled queries over a 69-document corpus, scoring retrieval and generation quality independently
- Ablation study: semantic-only retrieval outperformed hybrid BM25+semantic retrieval on every metric (Recall@5: 0.848 vs. 0.818, Faithfulness: 0.788 vs. hybrid), overturning the assumed hybrid-is-better default
- Eliminated multi-query retrieval after it matched or underperformed on every metric while approximately doubling LLM calls per query; evaluated k=3, k=5, and k=8, selecting k=5 as the optimal precision/recall/answer-relevancy tradeoff

Personal Intelligence System | Python · FastAPI · OpenRouter · Multi-Agent

Ongoing

- Built a self-hosted personal knowledge engine with 13 plug-and-play connectors (Notion, Slack, Email/IMAP, GitHub, Jira, Linear, Confluence, and more) that chunk, embed, and store all ingested content in a semantic vector space queryable in natural language
- Architected a two-layer agent system: Boss Agent auto-schedules syncs across all connected sources; Brain Agent dispatches to per-source sub-agents with retry/backoff and hash-based content dedup to skip unchanged items

SLM Pretraining From Scratch | PyTorch · Python

June – July 2026

- Implemented a decoder-only transformer (GPT-style architecture) from scratch, including attention mechanism, tokenizer, and training loop, without pretrained weights or external frameworks; 2.1M parameters
- Trained to a validation loss of 1.74 (~5.7 perplexity) using cosine learning rate decay with warmup; evaluated 5 architecture configurations from 120K to 4.8M parameters to characterize the relationship between model size and overfitting in a low-data regime

TrenTorch | Python · NumPy · OSS

Aug 2026 – Present

- Core engineer on TrenTorch (github.com/TrenTorch/TrenTorch) — open-source educational framework teaching PyTorch internals by rebuilding every component from scratch; own the tren CLI layer: test runner, Jupyter integration, progress tracking, benchmarking

EDUCATION

Vellore Institute of Technology

Vellore, India

Bachelor of Technology in Information Technology, 2024 – 2028

SKILLS

Languages: Python, SQL, C

RAG / Retrieval: LangChain, RAGAS, Chroma, FAISS, Embeddings, Vector Databases, Semantic Search, CRAG

Modeling: PyTorch, Machine Learning, Deep Learning, Transformer Architecture, Fine-Tuning (LoRA/PEFT), Prompt Engineering

Backend / Tools: Hugging Face, FastAPI, Flask, React, Pandas, NumPy, Docker, Git, PostgreSQL, MySQL